

Bayesian model assessment: Use of conditional vs marginal likelihoods

E. C. Merkle

University of Missouri

D. Furr and S. Rabe-Hesketh

University of California, Berkeley

Abstract

Typical Bayesian methods for models with latent variables (or random effects) involve directly sampling the latent variables along with the model parameters. In high-level software code for model definitions (using, e.g., BUGS, JAGS, Stan), the likelihood is therefore specified as conditional on the latent variables. This can lead researchers to perform model comparisons via conditional likelihoods, where the latent variables are considered model parameters. In other settings, typical model comparisons involve marginal likelihoods where the latent variables are integrated out. This distinction is often overlooked despite the fact that it can have a large impact on the comparisons of interest. In this paper, we clarify and illustrate these issues, focusing on the comparison of conditional and marginal Deviance Information Criteria (DICs) and Watanabe-Akaike Information Criteria (WAICs) in psychometric modeling. The conditional/marginal distinction corresponds to whether the model should be predictive for the clusters that are in the data or for new clusters (where “clusters” typically correspond to higher-level units like people or schools). Correspondingly, we show that marginal WAIC corresponds to leave-one-cluster out (LOcO) cross-validation, whereas conditional WAIC corresponds to leave-one-unit (LOuO). These results lead to recommendations on the general application of these criteria to models with latent variables.

Keywords: Bayesian information criteria, conditional likelihood, cross-validation, DIC, IRT, leave-one-cluster out, marginal likelihood, MCMC, SEM, WAIC.

1 Introduction

Psychometric models typically include latent variables (called “random effects” in some cases), classic examples being item response theory (IRT) models, structural equation models (SEMs) and multilevel/hierarchical regression models (MLMs). In IRT and SEM, the latent variables represent latent traits or characteristics of *persons*, and in MLMs, they represent random intercepts and possibly random slopes associated with *clusters*, such as schools, countries or, in the case of longitudinal data, persons. We will use the generic term “clusters” for the entities represented by latent variables in the model. In a Bayesian setting, latent variable models are hierarchical Bayesian models with latent variables as “direct” parameters. Bayesian approaches have recently received increased interest due to improvements in the ease of coding and estimating complex models by Markov chain Monte Carlo (MCMC) in various software packages. When it is computationally demanding to evaluate the *marginal* likelihood (over the latent variables) for likelihood-based inference, MCMC has the great advantage of working with the *conditional* likelihood (given the latent variables) by sampling the latent variables in tandem with the other model parameters.

The sampling of latent variables can lead to a decision point during model evaluation, when metrics such as the Deviance Information Criterion (DIC; Spiegelhalter, Best, Carlin, & van der Linde, 2002) or the Watanabe-Akaike Information Criterion (WAIC; Watanabe, 2010) are used to compare models. This decision point involves the question of whether or not the latent variables should be treated as model parameters. If so, then we would compute information criteria via conditional likelihoods, and we can expect “effective number of parameter” metrics to be large (each cluster has one or more unique latent variable values, and there are many clusters). If latent variables are not treated as model parameters, then we would compute information criteria via marginal likelihoods that are not generally available from the MCMC output.

Because DIC and WAIC are measures of models’ generalizability to new units, the “conditional vs. marginal” issue can be framed as a question of the type of data to which we wish to generalize. If we want our models to generalize to new units from the same clusters as in our original data, then the conditional likelihood (conditioning on the latent variables of the specific clusters observed) is appropriate. Conversely, if we want our models to generalize to new units from clusters not in the original data, then the marginal likelihood is appropriate. We can also think of this distinction in terms of how we would define the folds in k -fold cross-validation and leave-one-out approaches, or in other words, what we would leave out during estimation in order to evaluate predictive accuracy. To generalize to new units from existing clusters, we would leave out individual units from these clusters, and the predictive distribution for these units would exploit information from other units in the same clusters by conditioning on the cluster-specific latent variable values. To generalize to new units from new clusters, we would leave out intact clusters, and the predictive distribution would be marginal over the latent variables associated with the clusters. From this standpoint, the marginal likelihood may generally be preferred in psychometric modeling; we typically wish to generalize inferences beyond the specific clusters that were observed in our data.

Spiegelhalter et al. (2002) discuss this distinction in their original paper on the DIC, where they refer to the issue as “model focus” (see also Celeux, Forbes, Robert, Tittering-

ton, et al., 2006; Millar, 2009; Trevisani & Gelfand, 2003). If the parameters “in focus” include the latent variables, the likelihood becomes what we call the conditional likelihood; otherwise the likelihood is what we call the marginal likelihood. However, the distinction is hardly ever discussed in applications. Bayesian books targeted at applied psychometricians (e.g., Fox, 2010; Song & Lee, 2012; Kaplan, 2014; Levy & Mislevy, 2016) and papers on Bayesian model comparison in psychometrics (e.g., Kang, Cohen, & Sung, 2009; F. Li, Cohen, Kim, & Cho, 2009; Zhu & Stone, 2012) define either the conditional DIC or a generic DIC for non-hierarchical models, then apply the conditional DIC without discussing the implications of this choice or mentioning the marginal DIC. The WAIC has not yet been used much in psychometrics, exceptions being Luo and Al-Harbi (2017) and daSilva, Bazán, and Huggins-Manley (2018), who define the conditional version of WAIC for IRT without mentioning the marginal alternative, and Lu, Chow, and Loken (2017) who use the marginal version in factor analysis, without discussing the issue or mentioning the conditional alternative. Moving beyond psychometrics, Gelman, Hwang, and Vehtari (2014) point out that there is a choice to be made when defining the likelihood for Bayesian information criteria (including WAIC) between what we call conditional and marginal versions, but they use only the conditional version of WAIC in their “8-schools data” application and comparison to leave-one-out approaches. We are aware of a small number of contributions that discuss the merits of the marginal version of WAIC (L. Li, Qui, & Feng, 2016; Millar, 2018), but these treatments are confined to the special case of one response per latent variable (i.e., one unit per cluster). Further, Vehtari, Mononen, Tolvanen, Sivula, and Winther (2016) define a marginal version of WAIC for Gaussian latent variable models without contrasting it to the conditional version.

From an applied perspective, researchers estimating Bayesian models are often unaware of the distinction between marginal and conditional versions of information criteria, relying on default software behavior. Users of BUGS (D. Lunn, Jackson, Best, Thomas, & Spiegelhalter, 2012; D. J. Lunn, Thomas, Best, & Spiegelhalter, 2000), JAGS (Plummer, 2003), and Stan (Stan Development Team, 2014) typically obtain information criteria based on a conditional likelihood, whereas users of blavaan (Merkle & Rosseel, in press) and Mplus (e.g., Muthén & Asparouhov, 2012) obtain information criteria based on a marginal likelihood. Therefore, a first contribution of our paper is to emphasize the importance of being aware of the conditional/marginal distinction for the DIC and its implications. We also clarify other differences in the definition of the DIC implemented in various pieces of software. A second contribution is to propose a marginal version of WAIC for the prevalent case of multiple units per cluster and to contrast it with the conditional version. We specifically show that the marginal WAIC corresponds to leave-one-cluster out (LOcO) cross-validation, whereas the conditional WAIC corresponds to leave-one-unit (LOuO) cross-validation and has no rationale in the case of one unit per cluster. A third contribution is to propose an efficient version of adaptive quadrature for the case where the marginal likelihood does not have a closed form. The method can be performed as a postprocessing step with MCMC samples of the model parameters and latent variables as input. Our idea of using the *marginal* (with respect to the model parameters) posterior first and second moments of the latent variables in the adaptive quadrature approximation can easily be applied in Monte Carlo integration by using the corresponding normal density as importance distribution, and this would be an improvement over the method proposed by L. Li et al. (2016). A fourth contribution is

to provide two real examples, one using SEM and the other using IRT, where conditional and marginal information criteria can lead to very different substantive conclusions. For these examples, a very large number of Monte Carlo draws is needed to obtain acceptable Monte Carlo errors for the information criteria, especially in the conditional case. We make some practical recommendations for assessing convergence of the information criteria.

In the next section, we define the class of models considered and in Section 3 we define and contrast information criteria based on conditional and marginal likelihoods. Two examples based on real data are presented in Section 4, one using SEMs (Section 4.1) and the other IRTs (Section 4.2). We end with a discussion in Section 5. The supplementary material includes R code to estimate the IRTs and SEMs presented in this paper and obtain the information criteria – by performing adaptive quadrature integration for IRT and by using the recently extended R package *blavaan* for SEM.

2 Models and conditional versus marginal likelihoods

We consider models with continuous latent variables, such as IRTs, SEMs and MLMs. Responses y_{ij} are observed for units i belonging to clusters j , such as students i in schools j or items/indicators i for persons j , whereas the latent variables ζ_j are cluster-specific. The responses follow a generalized linear (or similar) model in which the conditional expectation of the responses is a function of the latent variables ζ_j and possibly observed covariates. The *conditional likelihood* is

$$f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) = \prod_{j=1}^J f_c(\mathbf{y}_j|\boldsymbol{\omega}, \boldsymbol{\zeta}_j), \quad (1)$$

where \mathbf{y} is the vector of responses for all units, \mathbf{y}_j is the vector of responses for cluster j ($j = 1, \dots, J$), $\boldsymbol{\zeta}$ is the vector of latent variables for all clusters, $\boldsymbol{\zeta}_j$ is the corresponding vector for cluster j , and $\boldsymbol{\omega}$ are model parameters. Under conditional independence, the joint conditional density (or probability) of the responses for cluster j factorizes as

$$f_c(\mathbf{y}_j|\boldsymbol{\omega}, \boldsymbol{\zeta}_j) = \prod_{i=1}^{n_j} f_c(y_{ij}|\boldsymbol{\omega}, \boldsymbol{\zeta}_j),$$

where $f_c(y_{ij}|\boldsymbol{\omega}, \boldsymbol{\zeta}_j)$ is the conditional density of the response for unit i in cluster j ($i = 1, \dots, n_j$), given $\boldsymbol{\omega}$, $\boldsymbol{\zeta}_j$, and possibly covariates (not shown).

For example, in a confirmatory factor analysis, y_{ij} is a continuous response for indicator i and person j , $\boldsymbol{\zeta}_j$ is a vector of common factors, $f_c(y_{ij}|\boldsymbol{\omega}, \boldsymbol{\zeta}_j)$ is a normal density, and $\boldsymbol{\omega}$ includes intercepts, factor loadings, and unique factor variances. In an item response model with binary items, $f(y_{ij}|\boldsymbol{\omega}, \boldsymbol{\zeta}_j)$ is $\pi_{ij}^{y_{ij}}(1 - \pi_{ij})^{1-y_{ij}}$ where π_{ij} is the conditional probability that y_{ij} equals 1 (often an inverse logit or inverse probit of a linear predictor), $\boldsymbol{\zeta}_j$ are dimensions of ability, and $\boldsymbol{\omega}$ includes difficulty parameters and possibly discrimination and guessing parameters. We call $f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})$ the conditional likelihood because it is conditional on the latent variables. Note that, in likelihood-based inference, “conditional likelihood” usually refers to the likelihood in which the latent variables have been eliminated by conditioning on sufficient statistics, whereas inference based on our conditional likelihood above is sometimes referred to as joint maximum likelihood estimation.

The latent variables are independent across clusters and have a joint prior distribution

$$g(\boldsymbol{\zeta}|\boldsymbol{\psi}) = \prod_{j=1}^J g(\zeta_j|\boldsymbol{\psi}_j)$$

with hyperparameters $\boldsymbol{\psi}$ that potentially include different values $\boldsymbol{\psi}_j$ for different (groups of) clusters j . Specification of the Bayesian model is completed by assuming prior distributions for $\boldsymbol{\omega}$ and $\boldsymbol{\psi}$, typically $p(\boldsymbol{\omega}, \boldsymbol{\psi}) = p(\boldsymbol{\omega})p(\boldsymbol{\psi})$.

The *marginal likelihood* is

$$f_m(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\psi}) = \prod_{j=1}^J \int f_c(\mathbf{y}_j|\boldsymbol{\omega}, \zeta_j) g(\zeta_j|\boldsymbol{\psi}_j) d\zeta_j. \quad (2)$$

We call this likelihood marginal because it is integrated over the latent variables. However, this likelihood is not marginal over the model parameters and should therefore not be confused with the marginal likelihood that is used, for instance for Bayes factors. The marginal likelihood defined above is the standard likelihood used in maximum likelihood estimation of latent variable or random-effects models. Celeux et al. (2006), who regard the latent variables as missing data, refer to this likelihood as the observed likelihood because it is marginal over the missing data.

Note that there are usually some parameters that can be included either in $\boldsymbol{\psi}$ or in $\boldsymbol{\omega}$, depending on how the model is formulated. For example, a linear two-level variance-components model can be written two alternative ways

$$\begin{aligned} \text{A: } & y_{ij} \sim N(\zeta_j, \sigma^2), \quad \zeta_j \sim N(\alpha, \tau^2) \\ \text{B: } & y_{ij} \sim N(\alpha + \zeta_j, \sigma^2), \quad \zeta_j \sim N(0, \tau^2) \end{aligned}$$

Version A corresponds to the two-stage formulation of hierarchical models (Raudenbush & Bryk, 2002), where latent variables (or random effects) appear in the model for y_{ij} and become “outcomes” in a second-stage model. In SEM, the first-stage and second-stage models are referred to as the measurement part and structural part, respectively. In the Bayesian literature, this formulation has been referred to as hierarchical centering (Gelfand, Sahu, & Carlin, 1995). If the model is formulated like this, the hyperparameters $\boldsymbol{\psi}$ include the mean α and variance τ^2 . Version B above specifies the reduced form model for y_{ij} so that ζ_j becomes a disturbance with zero mean and there is now only one hyperparameter, $\boldsymbol{\psi} = \tau^2$.

3 Conditional and marginal information criteria

To illustrate problems and differences associated with the use of conditional (vs marginal) likelihoods, we focus on two specific metrics: DIC and WAIC. The former is included due to its popularity (related to its inclusion in BUGS and JAGS), whereas the latter is included because it is closer to the current “start of the art” in Bayesian model assessment. We do not consider simpler criteria like AIC or BIC, but similar results can be expected. This is because AIC and BIC include parameter counts, which will change greatly depending on whether the marginal or conditional likelihood is used.

3.1 DIC

The deviance information criterion (DIC) was introduced by Spiegelhalter et al. (2002). For a model with parameters estimated as $\tilde{\boldsymbol{\theta}}$, they define the loss in assigning the model-implied *predictive* density $f(\mathbf{y}^r|\tilde{\boldsymbol{\theta}})$ to new, replicate (or out-of-sample) data \mathbf{y}^r as $-2\log f(\mathbf{y}^r|\tilde{\boldsymbol{\theta}})$. Following Efron (1986), they express the expectation of this out-of-sample loss over the distribution of the replicate data as the sum of the in-sample loss (evaluated for the observed data) and the “optimism” due to using the same data to estimate $\boldsymbol{\theta}$ and evaluate the loss:

$$\mathbb{E}_{\mathbf{y}^r}[-2\log f(\mathbf{y}^r|\tilde{\boldsymbol{\theta}})] = -2\log f(\mathbf{y}|\tilde{\boldsymbol{\theta}}) + \text{optimism}, \quad (3)$$

where the first term is sometimes referred to as the “plug-in deviance.” Spiegelhalter et al. (2002) use heuristic arguments to approximate the optimism by $2p_D$, where

$$p_D = \mathbb{E}_{\boldsymbol{\theta}|\mathbf{y}}[-2\log f(\mathbf{y}|\boldsymbol{\theta})] + 2\log f(\mathbf{y}|\tilde{\boldsymbol{\theta}}), \quad (4)$$

the posterior expectation of the deviance minus the deviance evaluated at $\tilde{\boldsymbol{\theta}}$, typically the posterior expectation of the parameters. Here p_D can be interpreted as a measure of the effective number of parameters. Combining these expressions, the DIC becomes

$$\begin{aligned} \text{DIC} &= -2\log f(\mathbf{y}|\tilde{\boldsymbol{\theta}}) + 2p_D \\ &= \mathbb{E}_{\boldsymbol{\theta}|\mathbf{y}}[-2\log f(\mathbf{y}|\boldsymbol{\theta})] + p_D \end{aligned} \quad (5)$$

Plummer (2008) proposes an alternative approximation for the optimism, which makes use of the undirected Kullback-Leibler divergence (also see Plummer’s contribution to the discussion of Spiegelhalter et al., 2002). In the general case, this alternative is given by

$$2p_D^P = \mathbb{E}_{\boldsymbol{\theta}|\mathbf{y}} \left[\log \left\{ \frac{f(\mathbf{y}^{r1}|\boldsymbol{\theta}^1)}{f(\mathbf{y}^{r1}|\boldsymbol{\theta}^2)} \right\} + \log \left\{ \frac{f(\mathbf{y}^{r2}|\boldsymbol{\theta}^2)}{f(\mathbf{y}^{r2}|\boldsymbol{\theta}^1)} \right\} \right], \quad (6)$$

where $\boldsymbol{\theta}^1$ and $\boldsymbol{\theta}^2$ are independent posterior draws (e.g., from parallel chains), \mathbf{y}^{r1} is a replicate drawn from $f(\mathbf{y}|\boldsymbol{\theta}^1)$, and \mathbf{y}^{r2} is defined similarly. For exponential family models, including the multivariate normal that is used later, the Kullback-Leibler divergence has a closed form so that replicates are not needed. JAGS uses this method (with replicates) for DIC computations, which is a difference between JAGS and BUGS that some researchers fail to realize. However, instead of substituting the above expression for the optimism in (3), or in other words for $2p_D$ in the first line of (5) as apparently recommended by Plummer (2008) and as implemented in blavaan, JAGS substitutes half this expression for p_D in the second line of (5). The JAGS approach is convenient as it does not require the postprocessing step of computing the plug-in deviance after the MCMC samples have been obtained.

Spiegelhalter et al. (2002) emphasize that, for hierarchical Bayesian models, the definition of p_D and DIC depends on which parameters are “in focus,” or in other words, which

parameters are included in θ . If the parameters in focus include the latent variables ζ , then $f(\cdot|\cdot)$ in equations (3) to (5) is $f_c(\mathbf{y}|\boldsymbol{\omega}, \zeta)$ from (1), whereas it is $f_m(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\psi})$ from (2) if the parameters in focus exclude the latent variables. We therefore define the conditional DIC as

$$\begin{aligned} \text{DIC}_c &= -2 \log f_c(\mathbf{y}|\tilde{\boldsymbol{\omega}}, \tilde{\zeta}) + 2p_{\text{Dc}} \\ &= E_{\boldsymbol{\omega}, \zeta|\mathbf{y}}[-2 \log f_c(\mathbf{y}|\boldsymbol{\omega}, \zeta)] + p_{\text{Dc}} \end{aligned} \quad (7)$$

and the marginal DIC as

$$\begin{aligned} \text{DIC}_m &= -2 \log f_m(\mathbf{y}|\tilde{\boldsymbol{\omega}}, \tilde{\boldsymbol{\psi}}) + 2p_{\text{Dm}} \\ &= E_{\boldsymbol{\omega}, \boldsymbol{\psi}|\mathbf{y}}[-2 \log f_m(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\psi})] + p_{\text{Dm}} \end{aligned} \quad (8)$$

with corresponding conditional and marginal variants of (4) or (6) for p_{Dc} and p_{Dm} . Note that Celeux et al.'s (2006) DIC_7 and DIC_1 correspond to our conditional and marginal DIC, respectively.

Conditioning on ζ implies that the new units \mathbf{y}^f in hypothetical replications come from the same clusters as the observed data, whereas marginalizing over ζ implies that the new units are for new clusters. We show in Appendix A that $E_{\boldsymbol{\omega}, \zeta|\mathbf{y}}[-2 \log f_c(\mathbf{y}|\boldsymbol{\omega}, \zeta)] \leq E_{\boldsymbol{\omega}, \boldsymbol{\psi}|\mathbf{y}}[-2 \log f_m(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\psi})]$ and in Appendix B that p_{Dc} tends to be much larger than p_{Dm} . Because the magnitudes of these differences will vary between models, there will be many instances where conditional DICs and marginal DICs favor different models. Plummer (2008) shows that the large-sample behavior of the penalty p_{D} depends on the dimension of θ and that the penalty may not approach the optimism if the dimensionality of θ increases with the sample size, as it does for latent variable models when ζ is in focus. This finding implies that the penalty term is a poor approximation of the optimism when the conditional DIC is used.

When the individual clusters are not of intrinsic interest or when we would like to choose among models that differ in the specification of the distribution for the latent variables $p(\zeta|\boldsymbol{\psi})$, then clearly the marginal DIC should be used. However, the conditional DIC is easier to compute because it does not require integration. For this reason, users of MCMC software (e.g., BUGS, JAGS) usually obtain the conditional DIC as output and are often unaware of the important distinction between conditional and marginal versions of the DIC.

It is sometimes possible to evaluate these marginal likelihoods by exploiting software for maximum likelihood estimation. For example, blavaan uses the related lavaan package (Rosseel, 2012) to evaluate the marginal likelihood for each MCMC iteration to obtain the posterior expectation of the deviance for p_{Dm} and once after MCMC sampling is complete to compute the plug-in deviance. In the case of normal likelihoods, as in linear mixed models or SEMs with continuous responses, the marginal likelihood has a closed form. However, closed-form solutions are generally not available for non-normal likelihoods, as in models with categorical responses. For such cases, we propose approximating the integrals using a computationally efficient form of adaptive quadrature. With posterior samples of the model

parameters and latent variables as input, the marginal deviances for each draw of the model parameters and for their posterior means are obtained by using the *marginal* (with respect to the model parameters) posterior first and second moments of the latent variables in the adaptive quadrature approximation. See Appendix C for details, where we also point out that the normal approximation to the marginal posterior distribution could be used as importance distribution in Monte-Carlo integration, which would be an improvement over standard Monte-Carlo integration used by L. Li et al. (2016).

3.2 WAIC

Gelman et al. (2014) and Vehtari, Gelman, and Gabry (2017) describe WAIC (Watanabe, 2010) as an approximation to minus twice the expected log pointwise predictive density (elpdd) for new data. The definition of the WAIC requires that the data points, such as responses y_{ij} for units, are independent given the parameters, so we first consider the conditional case where the parameters θ include the latent variables. In this case, minus twice elppd is

$$-2 \sum_{j=1}^J \sum_{i=1}^{n_j} E_{y_{ij}^r} \log \left[E_{\omega, \zeta | \mathbf{y}} f_c(y_{ij}^r | \omega, \zeta_j) \right] = -2 \sum_{j=1}^J \sum_{i=1}^{n_j} \log \left[E_{\omega, \zeta | \mathbf{y}} f_c(y_{ij} | \omega, \zeta_j) \right] + \text{optimism.} \quad (9)$$

where $E_{\omega, \zeta | \mathbf{y}} f_c(y_{ij}^r | \omega, \zeta_j)$ on the left is the pointwise (posterior) predictive density for a *new* response y_{ij}^r and its logarithm is averaged over the distribution of new responses. In contrast, the first term on the right is evaluated at the *observed* response y_{ij} and therefore represents the in-sample version of the elppd, called lppd by Gelman et al. (2014). Unlike the plug-in deviance in the DIC_c , where we condition on point estimates $\tilde{\omega}, \tilde{\zeta}$, here we take the expectation of $f_c(y_{ij}^r | \omega, \zeta_j)$ over the posterior distribution of ω and ζ . This can be written as

$$E_{\omega, \zeta | \mathbf{y}} f_c(y_{ij}^r | \omega, \zeta_j) = \int f_c(y_{ij}^r | \omega, \zeta_j) \left[\int p(\zeta_j | \mathbf{y}_j, \omega, \psi) p(\omega, \psi | \mathbf{y}) d\psi \right] d\omega d\zeta_j,$$

where the term in square brackets evaluates to the joint posterior $p(\omega, \zeta_j | \mathbf{y})$. It is clear from this expression that, even after conditioning on ω and ψ , whose posterior depends on all the data, the data \mathbf{y}_j for cluster j provides direct information on ζ_j . As discussed by Gelman, Meng, and Stern (1996), this *posterior predictive density* is appropriate for the situation where hypothetical new data are responses from the existing clusters. The conditional WAIC is then given by

$$\text{WAIC}_c = -2 \sum_{j=1}^J \sum_{i=1}^{n_j} \log \left[E_{\omega, \zeta | \mathbf{y}} f_c(y_{ij} | \omega, \zeta_j) \right] + 2p_{\text{Wc}}, \quad (10)$$

where p_{Wc} is the effective number of parameters and can be approximated by (see Gelman et al. (2014) for an alternative approximation)

$$p_{\text{Wc}} = \sum_{j=1}^J \sum_{i=1}^{n_j} \text{Var}_{\omega, \zeta | \mathbf{y}} [\log f_c(y_{ij} | \omega, \zeta_j)].$$

Vehtari et al. (2017) discuss problems with this approximation when any of the posterior variances of the log-pointwise predictive densities, $\text{Var}_{\omega, \zeta | \mathbf{y}}[\log f_c(y_{ij} | \omega, \zeta_j)]$, are too large, giving 0.4 as an upper limit based on simulation evidence.

If the hypothetical new data come from new clusters with new values ζ_j^r of the latent variables, the *mixed predictive density* (Gelman et al., 1996; Marshall & Spiegelhalter, 2007) should be used, which is the posterior expectation of $f_m(y_{ij}^r | \omega, \psi)$ and can be written as

$$E_{\omega, \psi | \mathbf{y}} f_m(y_{ij}^r | \omega, \psi) = \int \left[\int f_c(y_{ij}^r | \omega, \zeta_j) p(\zeta_j | \psi) d\zeta_j \right] p(\omega, \psi | \mathbf{y}) d\omega d\psi, \quad (11)$$

where the term in brackets, which evaluates to $f_m(y_{ij}^r | \omega, \psi)$, involves only the prior of ζ_j . When making predictions for a new cluster, the data provide information on ζ_j only indirectly by providing information on ψ . However, for clustered data, notice that the responses for different units are not conditionally independent given ω and ψ , so that we must redefine the “data points” as the vectors of responses for the clusters \mathbf{y}_j . The marginal WAIC now becomes

$$\text{WAIC}_m = -2 \sum_{j=1}^J \log \left[E_{\omega, \psi | \mathbf{y}} f_m(\mathbf{y}_j | \omega, \psi) \right] + 2p_{Wm}$$

with

$$p_{Wm} = \sum_{j=1}^J \text{Var}_{\omega, \psi | \mathbf{y}} [\log f_m(\mathbf{y}_j | \omega, \psi)],$$

where $f_m(\mathbf{y}_j | \omega, \psi)$ is given as term j of the product in (2). L. Li et al. (2016) and Millar (2018) define WAIC_m for the case when there is only one unit per cluster. Examples are overdispersed Poisson or binomial regression models, where the unit-level latent variable modifies the variance function, and meta-analysis, where the data for the clusters have been aggregated into effect-size estimates.

3.3 Interpretation of WAIC via connections with LOO-CV

Watanabe (2010) showed that WAIC and Bayesian leave-one out (LOO) cross-validation (CV) are asymptotically equivalent. Following our previous discussion, the different forms of WAIC represent different forms of LOO-CV. Because the predictive distribution in the conditional WAIC is for a response from a new unit in an existing cluster, it approximates leave one *unit* out CV, given as

$$-2\text{LOuO-CV} = -2 \sum_{j=1}^J \sum_{i=1}^{n_j} \log E_{\omega, \zeta | \mathbf{y}_{-ij}} f_c(y_{ij} | \omega, \zeta_j),$$

where \mathbf{y}_{-ij} represents the full data with observation ij held out. In contrast, the predictive distribution in the marginal WAIC is the joint distribution for the responses of all units in a new cluster and therefore approximates leave one *cluster* out CV:

$$-2\text{LOcO-CV} = -2 \sum_{j=1}^J \log E_{\omega, \psi | \mathbf{y}_{-j}} f_m(\mathbf{y}_j | \omega, \psi).$$

In other words, the marginal WAIC generalizes to new clusters whereas the conditional WAIC generalizes only to new units in the clusters that are in the data. Note that the first term in WAIC_m , the in-sample lppd, corresponds to the “hierarchical approximation” (Vehtari, Mononen, et al., 2016) or the “full-data mixed” (Marshall & Spiegelhalter, 2007) method for approximating what we call LOcO-CV.

L. Li et al. (2016) and Millar (2018) motivate their marginal WAIC_m as an approximation to LOO-CV in the case of one unit per cluster. Millar (2018) shows that conditional leave-one out (our LOuO) becomes marginal leave-one out (our LOcO) when there is only one unit per cluster, because there are no data for the cluster to condition on after removing the unit. Therefore, WAIC_c has no clear justification in this case. Vehtari et al. (2017) nevertheless use WAIC_c for the “eight schools” data, a meta-analysis of SAT preparatory programs. They find that WAIC_c diverges from LOO-CV when the responses are multiplied by a factor S with $S > 1.5$ (see their Figure 1a). We replicated their results and found that WAIC_m , which Vehtari et al. (2017) did not consider, continues to be a good approximation to LOO-CV for $S = 4$, for which we obtained $\text{LOO-CV} = 86.0$, $\text{WAIC}_m = 85.5$, and $\text{WAIC}_c = 68.7$. Millar (2018) also points out that two regularity conditions required for asymptotic equivalence of WAIC and LOO-CV do not hold for the conditional versions, namely (1) the observations are not identically distributed (their distribution depends on the latent variables ζ_j) and (2) the number of parameters increases with the sample size. These two regularity conditions are also violated in the clustered case, so it is not clear whether WAIC_c is a good approximation to LOuO-CV.

Vehtari et al. (2017) introduce a computationally efficient approximation to LOO using Pareto-smoothed importance sampling (PSIS) and implement this method in R package loo (Vehtari, Gelman, & Gabry, 2016). This package computes both PSIS-LOO and WAIC with MCMC draws of pointwise predictive densities as input, and it can be used to compute PSIS-LOuO, PSIS-LOcO, WAIC_c , and WAIC_m by providing either $\log f_c(y_{ij}|\boldsymbol{\omega}, \zeta_j)$ or $\log f_m(\mathbf{y}_j|\boldsymbol{\omega}, \boldsymbol{\psi})$ as input. This package is used in Section 4.2.

4 Examples

We now discuss two examples to highlight differences in various types of Bayesian information criteria. The first example involves a series of increasingly-constrained factor analysis models, as might be used in a measurement invariance study. This example focuses on DIC and its various definitions across software and across conditional/marginal likelihoods. The second example involves item response models, where marginal versions of the criteria do not have closed forms. In this example, we consider WAIC and PSIS-LOO in addition to DIC.

4.1 Measurement Invariance Study via Multi-Group CFA

Wicherts, Dolan, and Hessen (2005) compared a series of four-group, one-factor confirmatory models using data from a stereotype threat experiment. Here we replicate these model comparisons via both conditional DIC and marginal DIC (the original comparisons involved frequentist models). This allows us to illustrate the practical impact of using different DICs.

The data consist of Differential Aptitude Test scores from 295 Dutch students comprising both majority and minority Dutch ethnicities. The test contained three subscales (verbal ability, mathematical ability, and logic), where each subscale score was computed

as the number of items that a student answered correctly.

The original authors used a “stereotype threat” manipulation during the data collection: half of the students were primed about ethnic stereotypes related to intelligence tests, while the other half received no primes. Wicherts et al. (2005) conducted a measurement invariance study with the resulting data, to examine the impact of the stereotype threat manipulation on the students. This measurement invariance study involved four groups: two ethnicities (majority, minority) crossed with the stereotype threat manipulation (received, not received). The measurement model for subscale i ($i = 1, 2, 3$) and person j ($j = 1, \dots, 295$) in group g ($g = 1, 2, 3, 4$) is

$$y_{ijg} | \nu_{ig}, \lambda_{ig}, \sigma_{ig}, \zeta_{jg} \sim N(\nu_{ig} + \lambda_{ig}\zeta_{jg}, \sigma_{ig}^2),$$

where ν_{ig} is a group-specific intercept, λ_{ig} is a group-specific factor loading for subscale i (with the loadings for the first subscale set to 1 for identification, $\lambda_{1g} = 1$), ζ_{jg} is the common factor, and σ_{ig}^2 is the group-specific variance of the unique factor for subscale i . The parameters ν_{ig} , λ_{ig} , and σ_{ig}^2 are in ω and have the following prior distributions for $g = 1, 2, 3, 4$:

$$\begin{aligned} \nu_{1g}, \nu_{2g}, \nu_{3g} &\sim N(0, 1000) \\ \lambda_{2g}, \lambda_{3g} &\sim N(0, 100) \\ \sigma_{1g}^{-2}, \sigma_{2g}^{-2}, \sigma_{3g}^{-2} &\sim \text{Gamma}(1, .5) \end{aligned}$$

The prior distribution of the common factor, or structural part of the SEM is

$$\zeta_{jg} \sim N(\alpha_g, \tau_g^2)$$

with group-specific means α_g and variances τ_g^2 . These hyperparameters are in ψ and have the following priors:

$$\begin{aligned} \alpha_g &\sim N(0, 100) \\ \tau_g^{-2} &\sim \text{Gamma}(1, 1), \end{aligned}$$

with the exception that α_1 is fixed to zero for identification.

The measurement invariance study involved the comparison of nine versions of the model above. The overall strategy was to set increasingly more parameters equal across groups (i.e., across the g subscript), starting with factor loadings λ_{ig} (models 2 and 2a), then unique factor variances σ_{ig}^2 (models 3 and 3a), then common factor variances τ_g^2 (model 4; constant across conditions but allowed to differ between minority and non-minority groups), and then intercepts, ν_{ig} (models 5, 5a and 5b), and finally factor means α_g (model 6; constant across conditions but allowed to differ between minority and non-minority groups). The models without letters resulted from the previous model by constraining one type of parameter, whereas the model with the same number and a letter resulted from freeing some of these parameters again based on SEM fit criteria and modification indices. For example, model 5 (intercepts restricted across all four groups) is nested in models 5a (one intercepts freed) and 5b (another intercept freed), and model 6 (factor means constrained equal across conditions) is nested in model 5b. See Wicherts et al. (2005) for further detail on the models.


We conduct a Bayesian re-analysis of the nine Wicherts et al. (2005) models here, focusing on model comparison via DIC. We used JAGS to estimate each of the models ten times; for each of these estimations, we used three chains of 1,000 adaptation iterations, 4,000 burnin iterations, and 10,000 sampling iterations. For each estimation, we computed four DIC statistics via *blavaan*: the conditional and marginal versions using both the Spiegelhalter et al. (2002) definition in (4) and the Plummer (2008) definition in (6) for “effective number of parameters.” Chain convergence was assessed via the Gelman-Rubin statistic (Gelman & Rubin, 1992), with the model judged to converge if all parameters exhibited a Gelman-Rubin statistic less than 1.2. The threshold of 1.2 is somewhat larger than the typically-recommended threshold of 1.05, and it was used here because, in our experience with these models, it can be difficult to simultaneously achieve the threshold of 1.05 on all parameters in JAGS even when the sampling is going well (especially when treating latent variables as parameters). Nonetheless, Gelman-Rubin values above 1.05 were rare; in estimating all nine models a single time, we found 18 parameters (out of over 2,700 total) with Gelman-Rubin statistics above 1.05, none of which had values over 1.10.

Spiegelhalter et al.’s (2002) conditional and marginal DICs for ten estimations of each of the nine models are displayed in Figure 1. Here the scale of the y axis for the marginal case is a 10-fold magnification of that for the conditional case, showing that the Monte Carlo error is far greater for the conditional DICs. However, there is substantial Monte Carlo error also in the marginal case, making model comparison unreliable for some pairs of models (e.g., Models 3 and 2a). This suggests that accurate DICs require sampling for more iterations, as compared to what is required for obtaining accurate parameter estimates.

Especially because these models were part of a measurement invariance study, we also examine where we would stop as we moved from Model 2 to 6 in a stepwise fashion (as was done in the original paper). In the marginal case, the DIC for Model 3 is generally larger than the DIC for Model 2a, so we may stop at Model 2a. However, Monte Carlo error is large enough that we may prefer Model 3 to Model 2a in a given replication, in which case we may stop at Model 4. But the marginal DIC is lowest for the final Model 6. The results are very different in the conditional case. Here, we would stop at Model 2a across all ten replications. Further, the models generally obtain larger DIC values as we move towards the final model 6, in stark contrast to the marginal case. Models 5b and 6 have the worst values in the conditional case, whereas they are the best in the marginal case.

Beyond the comparison of DICs, we further examine the four definitions of effective number of parameters: marginal versus conditional crossed with the Spiegelhalter versus Plummer definitions. These are displayed for the nine Wicherts et al. (2005) models in Figure 2, with the lines showing simple parameter counts for each model (where, in the conditional case, latent variables count as parameters). Results are similar to the previous figure in that the conditional definitions exhibit greater Monte Carlo error than the marginal definitions (note the differences in y -axis scales, both within this figure and as compared to Figure 1). In the marginal case, we are counting the unrestricted item-specific parameters in ω and the unrestricted factor means and variances in ψ (note that these parameter counts increase when parameters are freed, i.e. going from model 2 to 2a, 3 to 3a, and from 5 to 5a to 5b.); the Plummer effective number of parameters is generally larger than the Spiegelhalter effective number of parameters here. In the conditional case, we are

Figure 1. Marginal and conditional DICs (Spiegelhalter et al. definitions) for nine models from Wicherts et al., where each model was estimated ten times.



counting the unrestricted item-specific parameters in ω and the common factors values for all subjects. Here we see greater discrepancy between the Spiegelhalter and Plummer definitions. The Spiegelhalter definition appears to work better here, as the values are nearly always smaller than the associated parameter counts (which we would expect based on the fact that the latent variables exhibit shrinkage). However, in fairness, Plummer (2008) explicitly states that his definition is not intended for situations where the number of parameters increase with the number of observations, as they do in the conditional case. Beyond this, we observe situations where the marginal effective number of parameters decreases while the conditional effective number of parameters increases or stays the same; the sequence from Models 4 to 6 exhibits multiple examples here.

Finally, we point out that much of the variability in DIC is due to Monte Carlo error in the estimate of the effective number of parameters. As shown in Appendix D, this Monte Carlo error is easy to estimate (in the Plummer case by ignoring any Monte Carlo error in the plug-in deviance). Figure 3 plots the Plummer DICs across the 10 replications, with error bars representing ± 2 times the estimated Monte Carlo error (from a single replication) only in the effective number of parameters. We observe that the error bars reflect the magnitude of variability in the overall DIC values. To support this claim, we computed the proportion of the time that the error bars covered the DICs across the ten replications and nine models. This coverage was 76% in the conditional case and 70% in the marginal case (the coverage for the corresponding effective numbers of parameters was 88% and 99%, respectively). While the coverage for the DIC is lower than one may hope, it illustrates that the deviance evaluated at the posterior means exhibits small variability compared to the variability in expected deviance (Spiegelhalter case) or K-L distance (Plummer case).

The results in this section illustrate that the use of marginal vs conditional DIC can impact substantive conclusions, in that different models can be selected with different versions of DIC and different conclusions can be reached about the effective number of pa-

Figure 2. Estimated effective number of parameters under marginal and conditional DIC for nine models from Wicherts et al., where each model was estimated ten times. Lines reflect simple parameter counts for each model, where latent variables count as parameters in the conditional case.

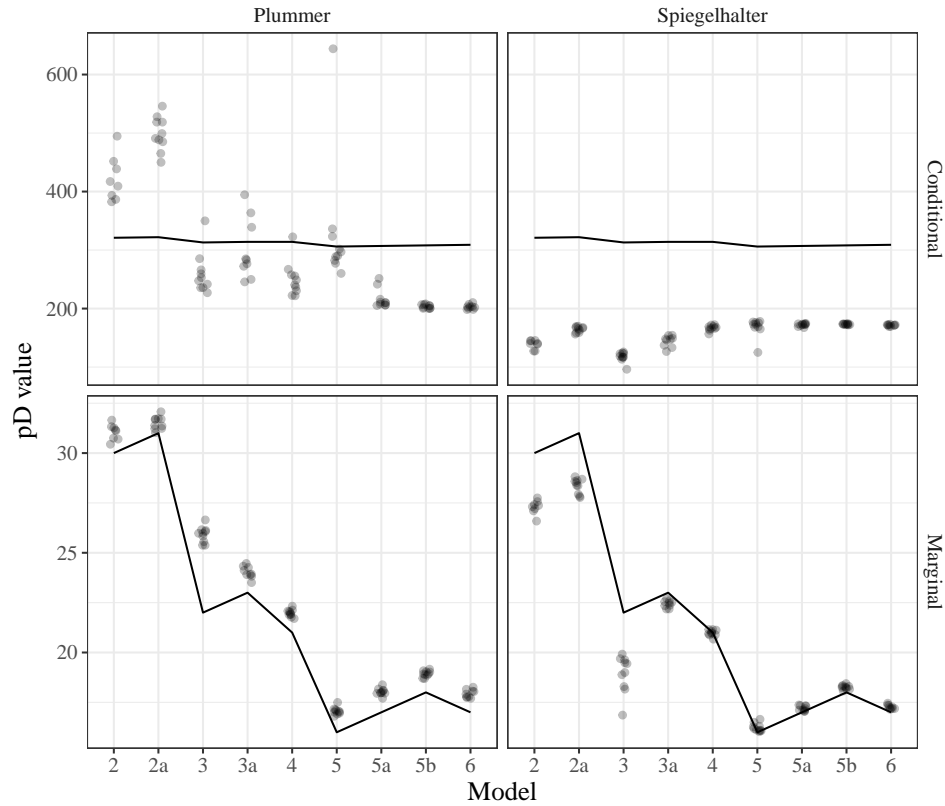



Figure 3. Marginal and conditional DICs (Plummer definitions) for ten replications of the Wicherts et al. models, with error bars (± 2 SDs) stemming from a single replication.



rameters. Furthermore, the conditional DIC exhibits larger Monte Carlo error and depends more strongly on the approximation use (Spiegelhalter versus Plummer) than the marginal version.

4.2 Explaining Individual Differences in Verbal Aggression via IRT

Here, we consider item response models with different person covariates to compare conditional and marginal versions of the information criteria. In the case of IRT, the marginal information criteria utilize the numerical integration methods described in Appendix C.

Several latent regression Rasch models are fit to the dataset on verbal aggression (Vansteelandt, 2000) that consists of $J = 316$ persons and $I = 24$ items. Participants were instructed to imagine four frustrating scenarios, and for each they responded to items regarding whether they would react by cursing, scolding, and shouting. They also responded to parallel items regarding whether they would *want* to engage in the three behaviors, resulting in six items per scenario (cursing/scolding/shouting \times doing/wanting). An example item is, “A bus fails to stop for me. I would want to curse.” The response options for all items were “yes”, “perhaps”, and “no.” The items have been dichotomized for this example by combining “yes” and “perhaps” responses. Two person-specific covariates are included: the respondent’s trait anger score (Spielberger, 1988), which is a raw score from a separate measure taking values between 11 to 39 in the data, and an indicator for whether the respondent is male, which takes the values 0 and 1.

The model that defines the conditional likelihood $f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})$ is

$$y_{ij}|\mathbf{x}_j, \boldsymbol{\gamma}, \zeta_j, \delta_i \sim \text{Bernoulli}\left(\text{logit}^{-1}(\mathbf{x}_j'\boldsymbol{\gamma} + \zeta_j - \delta_i)\right), \quad (12)$$

where $y_{ij} = 1$ if person j responds positively to item i and $y_{ij} = 0$ otherwise, \mathbf{x}_j is a vector of a constant and $K - 1$ person-specific covariates, $\boldsymbol{\gamma}$ is a vector of regression coefficients, ζ_j is a person residual, and δ_i is an item difficulty parameter. The last item difficulty is constrained, $\delta_I = -\sum_{i=1}^{(I-1)} \delta_i$. The priors for the item difficulties and regression coefficients that comprise $\boldsymbol{\omega}$ are

$$\begin{aligned} \delta_1, \dots, \delta_{I-1} &\sim \text{N}(0, 9) \\ \gamma_1^*, \dots, \gamma_K^* &\sim t_1(0, 1). \end{aligned}$$

Unlike the previous example, ζ_j is just a disturbance here with zero mean and hyperparameter $\psi = \tau$:

$$\begin{aligned} \zeta_j &\sim \text{N}(0, \tau^2) \\ \tau &\sim \text{Exp}(.1). \end{aligned}$$

Weakly informative priors are specified for the regression coefficients $\boldsymbol{\gamma}$, following Gelman, Jakulin, Pittau, and Su (2008). Specifically, after standardizing continuous covariates to have zero mean and standard deviation 0.5 and binary covariates to have mean 0 and

range 1, t -distributions with one degree of freedom are used as priors for the corresponding regression coefficients, γ^* . The MCMC draws of the coefficients are then transformed to reflect the original scales of the covariates.

Focus is placed on ζ_j for the conditional approach, which yields a prediction inference involving new responses from the same persons (and items). The marginal approach, perhaps more realistically, places focus on τ , implying a prediction inference involving new responses from a new persons. Five competing models are considered, differing only in what person covariates are included: Model 1 includes no covariates ($K = 1$), Model 2 has the trait anger score ($K = 2$), Model 3 has the indicator for male ($K = 2$), Model 4 has both covariates ($K = 3$), and Model 5 has both covariates and their interaction ($K = 4$).


The five models are estimated with Stan using 5 chains of 2,500 draws with the first 500 draws of each discarded, resulting in a total of 10,000 kept posterior draws. The unusually large number of posterior draws is chosen here due to the anticipated Monte Carlo errors for the information criteria, but such a large number is not ordinarily necessary for estimating the posterior means and standard deviations of the parameters. The adaptive quadrature method described in Appendix C is used to integrate out ζ for computing the marginal DIC and WAIC and PSIS-LOO. Eleven quadrature points were found to be sufficient for one of the models (Model 4) using the method described in Appendix C, and this number of points is used throughout.

Figure 4 provides the estimates for the information criteria and PSIS-LOO, where the conditional and marginal versions now have the same y -axis units. The DIC and WAIC differ from each other for any given model, though they seem to show a similar pattern between models. The WAIC is a much better approximation to the PSIS-LOO in the marginal case than the conditional case. The high degree of Monte Carlo error in the conditional versions renders differentiating the predictive performance of the models difficult. In the marginal case, the amount of Monte Carlo error is less but still poses a degree of difficulty in making comparisons, but Models 1 and 3 now clearly provide poorer predictions in comparison to the others. These models do not include trait anger, suggesting that this variable is an important predictor of verbal aggression. There is some evidence supporting Model 4 (without the interaction) as the best among the candidates.

Figure 5 provides the estimated effective number of parameters (dots) and actual number of parameters (lines) for the conditional and marginal versions of the DIC and WAIC. For conditional information criteria, the effective number of parameters are substantially less than the numbers of parameters, owing mainly to the fact that each ζ_j , with its hierarchical prior, contributes less than one to the effective number of parameters. For marginal information criteria, on the other hand, the effective number of parameters are close to the numbers of model parameters and somewhat larger than the number of parameters for WAIC.

For the WAIC, contributions to the effective number of parameters (from item-person combinations for p_{Wc} and from persons for p_{Wm}) should be less than .4 which was always satisfied for the marginal WAIC but was violated for an average of between 2.1 and 3.4 contributions for the conditional WAIC across the five models (for details, see Furr, 2017).

Figure 4. Information criteria for the five latent regression Rasch models. Points represent the results of the 10 independent MCMC simulations per model. A small amount of horizontal jitter is added to the points.




5 Discussion

In this paper, we sought to clarify and illustrate the various forms of Bayesian information criteria that are used in practice. Researchers often rely on pre-packaged software to obtain these criteria, failing to realize the fact that multiple versions of the criteria are available depending on the software (BUGS vs. JAGS) and depending on whether or not the latent variables appear in the model likelihood (leading to *conditional* and *marginal* information criteria, respectively). Additionally, even if researchers are aware of these distinctions, we showed that the criteria can suffer from large amounts of Monte Carlo error. Hence, long chains will often be necessary to obtain precise point estimates of the criteria. If ignored, these issues can lead to model comparisons that are irrelevant and/or irreproducible.

Similarly to the use of marginal likelihoods in general latent variable models, we recommend use of marginal information criteria as a default in Bayesian analyses. As previously discussed, marginal criteria assess a model's predictive ability when applied to new clusters. This is exactly what is desired in many psychometric contexts, where clusters may be defined by, e.g., countries, schools, or students. In these cases, we wish to discern general properties of scales or items that are not specific to the clusters that were observed. As a side advantage, the marginal information criteria also tend to have less Monte Carlo error than their conditional counterparts.

There also appear to be other problems in the case of conditional information criteria. For the DIC, Plummer (2008) showed that the penalty does not approach the optimism when the number of parameters increase with the sample size (as they do in the conditional case) and we observed in the first example that the penalty term depended largely on the

Figure 5. Estimated effective number of parameters (\hat{p}) for the five latent regression Rasch models. Points represent the results of the 10 independent MCMC simulations per model. A small amount of horizontal jitter is added to the points. The horizontal lines represent the counts of parameters associated with each model and focus.



definition used (Plummer versus Spiegelhalter). In the second example, we observed that the WAIC differed much more from PSIS-LOO in the conditional case than in the marginal case. This corresponds to the work of Millar (2018), who points out that two regularity conditions are not satisfied in the conditional case, one of them again being the number of parameters increasing with sample size.

Because the chain length necessary to obtain precise information criteria is generally larger than the length required to obtain stable posterior means of individual parameters, it will often be useful to compute the information criteria at multiple points. First, the criteria can be computed when individual parameters' posterior means have converged and stabilized. The researcher can then continue running the chains for the same number of iterations, re-computing the information criteria and effective number of parameters, and assessing changes in results. This process can be continued until the information criteria have reached the desired level of precision. Our examples suggest that the Monte Carlo error of the effective number of parameters could also be used as a lower bound of the imprecision of the information criteria.

Finally, while the dichotomous “conditional/marginal” distinction is most relevant for psychometric models, middle grounds are possible in some situations. For example, in an item response context, De Boeck (2008) proposes modeling both the item parameters and person parameters as random effects. Here, a researcher might be interested in the model's predictive ability when new people complete the same items that originally entered in the model. We would then compute information criteria where the person random effects are marginalized out of the likelihood but the item random effects are not. These types of

scenarios become relevant in datasets with multiple clustering variables, so that the models may have more than two “levels” and/or (partially-)crossed random effects.

6 References

- Celeux, G., Forbes, F., Robert, C. P., Titterton, D. M., et al. (2006). Deviance information criteria for missing data models. *Bayesian Analysis*, 1(4), 651–673.
- daSilva, M. A., Bazán, J. L., & Huggins-Manley, A. C. (2018). Sensitivity analysis and choosing between alternative polytomous irt models using Bayesian model comparison criteria. *Communications in Statistics - Simulation and Computation*, in press.
- De Boeck, P. (2008). Random item IRT models. *Psychometrika*, 73, 533–559.
- Fox, J.-P. (2010). *Bayesian item response modeling: Theory and applications*. NY: Springer.
- Furr, D. C. (2017). *Bayesian and frequentist cross-validation methods for explanatory item response models* (Unpublished doctoral dissertation). University of California, Berkeley, CA.
- Gelfand, A. E., Sahu, S. K., & Carlin, B. P. (1995). Efficient parametrisations for normal linear mixed models. *Biometrika*, 82, 379–488.
- Gelman, A., Hwang, J., & Vehtari, A. (2014). Understanding predictive information criteria for Bayesian models. *Statistics and Computing*, 24, 997–1016.
- Gelman, A., Jakulin, A., Pittau, M. G., & Su, Y.-S. (2008). A weakly informative default prior distribution for logistic and other regression models. *The Annals of Applied Statistics*, 1360–1383.
- Gelman, A., Meng, X. L., & Stern, H. (1996). Posterior predictive assessment of model fitness via realized discrepancies. *Statistica Sinica*, 6, 733–807.
- Gelman, A., & Rubin, D. B. (1992). Inference from iterative simulation using multiple sequences (with discussion). *Statistical Science*, 7, 457–511.
- Kang, T., Cohen, A. S., & Sung, H.-J. (2009). Model selection indices for polytomous items. *Applied Psychological Medicine*, 35, 499–518.
- Kaplan, D. (2014). *Bayesian statistics for the social sciences*. NY: The Guildford Press.
- Levy, R., & Mislevy, R. J. (2016). *Bayesian psychometric modeling*. Boca Raton, FL: Chapman & Hall.
- Li, F., Cohen, A. S., Kim, S. H., & Cho, S. J. (2009). Model selection methods for mixture dichotomous IRT models. *Applied Psychological Measurement*, 33, 353–373.
- Li, L., Qui, S., & Feng, C. X. (2016). Approximating cross-validatory predictive evaluation in Bayesian latent variable models with integrated IS and WAIC. *Statistics and Computing*, 26, 881–897.
- Lu, Z.-H., Chow, S.-M., & Loken, E. (2017). A comparison of Bayesian and frequentist model selection methods for factor analysis models. *Psychological Methods*, 22(2), 361–381.
- Lunn, D., Jackson, C., Best, N., Thomas, A., & Spiegelhalter, D. (2012). *The BUGS book: A practical introduction to Bayesian analysis*. Chapman & Hall/CRC.
- Lunn, D. J., Thomas, A., Best, N., & Spiegelhalter, D. (2000). WinBUGS – a Bayesian modelling framework: Concepts, structure, and extensibility. *Statistics and Computing*, 10, 325–337.
- Luo, U., & Al-Harbi, K. (2017). Performances of LOO and WAIC as IRT model selection methods. *Psychological Test and Assessment Modeling*, 59, 183–205.
- Marshall, E. C., & Spiegelhalter, D. J. (2007). Identifying outliers in Bayesian hierarchical models: A simulation-based approach. *Bayesian Analysis*, 2(2), 409–444.
- Merkle, E. C., & Rosseel, Y. (in press). *blavaan*: Bayesian structural equation models via parameter expansion. *Journal of Statistical Software*.
- Millar, R. B. (2009). Comparison of hierarchical Bayesian models for overdispersed count data using DIC and Bayes’ factors. *Biometrics*, 65, 962–969.

- Millar, R. B. (2018). Conditional vs. marginal estimation of predictive loss of hierarchical models using waic and cross-validation. *Statistics and Computing*, *in press*.
- Muthén, B., & Asparouhov, T. (2012). Bayesian structural equation modeling: A more flexible representation of substantive theory. *Psychological Methods*, *17*, 313–335.
- Pinheiro, J. C., & Bates, D. M. (1995). Approximations to the log-likelihood function in the nonlinear mixed-effects model. *Journal of Computational Graphics and Statistics*, *4*, 12–35.
- Plummer, M. (2003). JAGS: A program for analysis of Bayesian graphical models using Gibbs sampling. In K. Hornik, F. Leisch, & A. Zeileis (Eds.), *Proceedings of the 3rd international workshop on distributed statistical computing*.
- Plummer, M. (2008). Penalized loss functions for Bayesian model comparison. *Biostatistics*, *9*(3), 523–539.
- Rabe-Hesketh, S., Skrondal, A., & Pickles, A. (2005). Maximum likelihood estimation of limited and discrete dependent variable models with nested random effects. *Journal of Econometrics*, *128*(2), 301–323.
- Raudenbush, S. W., & Bryk, A. S. (2002). *Hierarchical linear models: Applications and data analysis methods* (2nd ed.). Thousand Oaks, CA: Sage.
- Rosseel, Y. (2012). *lavaan*: An R package for structural equation modeling. *Journal of Statistical Software*, *48*(2), 1–36.
- Song, X.-Y., & Lee, S.-Y. (2012). *Basic and advanced Bayesian structural equation modeling: With applications in the medical and behavioral sciences*. Chichester, UK: Wiley.
- Spiegelhalter, D. J., Best, N. G., Carlin, B. P., & van der Linde, A. (2002). Bayesian measures of model complexity and fit. *Journal of the Royal Statistical Society Series B*, *64*, 583–639.
- Spielberger, C. (1988). State-trait anger expression inventory research edition [Computer software manual]. Odessa, FL.
- Stan Development Team. (2014). Stan modeling language users guide and reference manual, version 2.5.0 [Computer software manual]. Retrieved from <http://mc-stan.org/>
- Trevisani, M., & Gelfand, A. E. (2003). Inequalities between expected marginal log-likelihoods, with implications for likelihood-based model complexity and comparison measures. *The Canadian Journal of Statistics*, *31*, 239–250.
- Vansteelandt, K. (2000). *Formal models for contextualized personality psychology* (Unpublished doctoral dissertation). University of Leuven, Leuven, Belgium.
- Vehtari, A., Gelman, A., & Gabry, J. (2016). loo: Efficient leave-one-out cross-validation and WAIC for Bayesian models. R package version 0.1.6. <https://github.com/stan-dev/loo>.
- Vehtari, A., Gelman, A., & Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, *27*, 1413–1432.
- Vehtari, A., Mononen, T., Tolvanen, V., Sivula, T., & Winther, O. (2016). Bayesian leave-one-out cross-validation approximations for Gaussian latent variable models. *Journal of Machine Learning Research*, 1–38.
- Watanabe, S. (2010). Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *Journal of Machine Learning Research*, *11*, 3571–3594.
- White, I. R. (2010). simsum: Analyses of simulation studies including Monte Carlo error. *The Stata Journal*, *10*, 369–385.
- Wicherts, J. M., Dolan, C. V., & Hessen, D. J. (2005). Stereotype threat and group differences in test performance: A question of measurement invariance. *Journal of Personality and Social Psychology*, *89*(5), 696–716.
- Zhu, X., & Stone, C. A. (2012). Bayesian comparison of alternative graded response models for performance assessment applications. *Educational and Psychological Measurement*, *72*(5), 774–799.

Appendix A

Posterior expectations of marginal and conditional likelihoods

Following Trevisani and Gelfand (2003), who studied DIC in the context of linear mixed models, we can use Jensen's inequality to show that the posterior expected value of the marginal log-likelihood is less than the posterior expected value of the conditional log-likelihood.

First, consider the function $h(x) = x \log(x)$. It is convex, so Jensen's inequality states that

$$h(E(x)) \leq E(h(x)). \quad (13)$$

Setting $x = f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})$ and taking expected values with respect to $\boldsymbol{\zeta}$, we have that

$$h(E(x)) = \log \left[\int f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \right] \int f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \quad (14)$$

$$E(h(x)) = \int \log(f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})) f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta}, \quad (15)$$

so that

$$\log \left[\int f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \right] \int f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \leq \int \log(f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})) f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta}. \quad (16)$$

We now multiply both sides of this inequality by $p(\boldsymbol{\omega}, \boldsymbol{\psi})/c$, where $p(\boldsymbol{\omega}, \boldsymbol{\psi})$ is a prior distribution and

$$c = \int_{\boldsymbol{\psi}} \int_{\boldsymbol{\omega}} \int_{\boldsymbol{\zeta}} f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \cdot p(\boldsymbol{\omega}, \boldsymbol{\psi}) d\boldsymbol{\omega} d\boldsymbol{\psi} \quad (17)$$

is the posterior normalizing constant. Finally, we integrate both sides with respect to $\boldsymbol{\omega}$ and $\boldsymbol{\psi}$ to obtain

$$\begin{aligned} \int_{\boldsymbol{\psi}} \int_{\boldsymbol{\omega}} \log \left(\int_{\boldsymbol{\zeta}} f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \right) \int_{\boldsymbol{\zeta}} f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \cdot [p(\boldsymbol{\omega}, \boldsymbol{\psi})/c] d\boldsymbol{\omega} d\boldsymbol{\psi} \leq \\ \int_{\boldsymbol{\psi}} \int_{\boldsymbol{\omega}} \int_{\boldsymbol{\zeta}} \log(f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})) f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta}) g(\boldsymbol{\zeta}|\boldsymbol{\psi}) d\boldsymbol{\zeta} \cdot [p(\boldsymbol{\omega}, \boldsymbol{\psi})/c] d\boldsymbol{\omega} d\boldsymbol{\psi}. \end{aligned} \quad (18)$$

We can now recognize both sides of (18) as expected values of log-likelihoods with respect to the model's posterior distribution, leading to

$$E_{\boldsymbol{\omega}, \boldsymbol{\psi}|\mathbf{y}} [\log f_m(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\psi})] \leq E_{\boldsymbol{\omega}, \boldsymbol{\zeta}|\mathbf{y}} [\log f_c(\mathbf{y}|\boldsymbol{\omega}, \boldsymbol{\zeta})]. \quad (19)$$

Note that the above results do not rely on normality, so they also apply to, e.g., the two-parameter logistic model estimated via marginal likelihood.

Appendix B

Effective number of parameters for marginal and conditional DIC

To consider the effective number of parameters for normal likelihoods, we rely on results from Spiegelhalter et al. (2002). They showed that the effective number of parameters p_D can be viewed as the fraction of information about model parameters in the likelihood, relative to the total information contained in both the likelihood and prior. Under this view, a specific model parameter gets a value of “1” if all of its information is contained in the likelihood, and it gets a value below “1” if some information is contained in the prior. We sum these values across all parameters to obtain p_D .

Spiegelhalter et al. (2002) relatedly showed that, for normal likelihoods, p_D can be approximated by

$$p_D \approx \text{tr}(\mathbf{I}(\hat{\boldsymbol{\theta}})\mathbf{V}), \quad (20)$$

where $\boldsymbol{\theta}$ includes all model parameters (including $\boldsymbol{\zeta}$ in the conditional model), $\mathbf{I}(\hat{\boldsymbol{\theta}})$ is the observed Fisher information matrix, and \mathbf{V} is the posterior covariance matrix of $\boldsymbol{\theta}$. When the prior distribution of $\boldsymbol{\theta}$ is noninformative, then $\mathbf{I}(\hat{\boldsymbol{\theta}}) \approx \mathbf{V}^{-1}$. Consequently, matching the discussion in the previous paragraph, the effective number of parameters under noninformative priors will approximate the total number of model parameters.

This result implies that the conditional p_D will tend to be much larger than the marginal p_D . In particular, in the conditional case, each individual has a unique $\boldsymbol{\zeta}$ vector that is included as part of the total parameter count. The resulting p_D will not necessarily be close to the total parameter count because the “prior distribution” of $\boldsymbol{\zeta}$ is a hyperdistribution, whereby individuals’ $\boldsymbol{\zeta}$ estimates are shrunk towards the mean. Thus, for these parameters, the “prior” is informative. However, even when the fraction of information in the likelihood is low for these parameters, the fact that we are summing over hundreds or thousands of $\boldsymbol{\zeta}$ vectors implies that the conditional p_D will be larger than the marginal p_D .

Appendix C

Adaptive Gaussian quadrature for marginal likelihoods

We modify the adaptive quadrature method proposed by Rabe-Hesketh, Skrondal, and Pickles (2005) for generalized linear mixed models to a form designed to exploit MCMC draws from the joint posterior of all latent variables and model parameters. Here we describe one-dimensional integration, but the method is straightforward to generalize to multidimensional integration as in Rabe-Hesketh et al. (2005). We assume that ζ_j represents a disturbance with zero mean and variance τ^2 so that there is only one hyperparameter $\psi = \tau$.

In a non-Bayesian setting, standard (non-adaptive) Gauss-Hermite quadrature can be viewed as approximating the conditional prior density $g(\zeta_j|\tau)$ by a discrete distribution with masses w_m , $m = 1, \dots, M$ at locations $a_m\tau$ so that the integrals in (2) are approximated by sums of M terms, where M is the number of quadrature points,

$$f_m(\mathbf{y}|\boldsymbol{\omega}, \tau) \approx \prod_{j=1}^J \sum_{m=1}^M w_m f_c(\mathbf{y}_j|\boldsymbol{\omega}, \zeta_j = a_m\tau). \quad (21)$$

To obtain information criteria, this method can easily be applied to the conditional likelihood function for each draw $\boldsymbol{\omega}^s$ and τ^s ($s = 1, \dots, S$) of the model parameters from MCMC output. This approximation can be thought of as a deterministic version of Monte Carlo integration. Adaptive quadrature is then a deterministic version of Monte Carlo integration via importance sampling.

If applied to each MCMC draw $\boldsymbol{\omega}^s$ and τ^s , the importance density is a normal approximation to the *conditional* posterior density of ζ_j , given the current draws of the model parameters. Rabe-Hesketh et al. (2005) used a normal density with mean and variance equal to the mean $E(\zeta_j|\mathbf{y}_j, \boldsymbol{\omega}^s, \tau^s)$ and variance $\text{var}(\zeta_j|\mathbf{y}_j, \boldsymbol{\omega}^s, \tau^s)$ of the conditional posterior density of ζ_j , whereas Pinheiro and Bates (1995) and some software use a normal density with mean equal to the mode of the conditional posterior and variance equal to minus the inverse Hessian of the conditional log posterior. Here, we modify the method by Rabe-Hesketh et al. (2005) for the Bayesian setting by using a normal approximation to the *unconditional* posterior density of ζ_j as importance density. Specifically, we use a normal density with mean

$$\tilde{\mu}_j = \tilde{E}(\zeta_j|\mathbf{y}) = \frac{1}{S} \sum_{s=1}^S \zeta_j^s, \quad (22)$$

and standard deviation

$$\tilde{\phi}_j = \sqrt{\widetilde{\text{var}}(\zeta_j|\mathbf{y}_j)} = \sqrt{\frac{1}{S-1} \sum_{s=1}^S (\zeta_j^s - \tilde{\mu}_j)^2}, \quad (23)$$

where ζ_j^s is the draw of ζ_j from its unconditional posterior in the s th MCMC iteration. The tildes indicate that these quantities are subject to Monte Carlo error.

Note that this version of adaptive quadrature is computationally more efficient than the one based on the mean and standard deviation of the *conditional* posterior distributions because the latter would have to be evaluated for each MCMC draw and would require numerical integration, necessitating a procedure that iterates between updating the quadrature locations and weights and updating the conditional posterior means and standard deviations. A disadvantage is that the quadrature locations and weights (and hence importance density) are not as targeted, but the computational efficiency gained also makes it more feasible to increase M .

The adaptive quadrature approximation to the marginal likelihood for cluster j at posterior draw s becomes

$$f_m(\mathbf{y}|\boldsymbol{\omega}, \tau) \approx \prod_{j=1}^J \sum_{m=1}^M w_{jm}^s f_c(\mathbf{y}_j|\boldsymbol{\omega}, \zeta_j = a_{jm}\tau^s), \quad (24)$$

where the adapted locations are

$$a_{jm} = \tilde{\mu}_j + \tilde{\phi}_j \times a_m, \quad (25)$$

and the corresponding weights or masses are

$$w_{jm}^s = \sqrt{2\pi} \times \tilde{\phi}_j \times \exp\left(\frac{a_{jm}^2}{2}\right) \times g(a_{jm}; 0, \tau^{2,s}) \times w_m, \quad (26)$$

where $g(a_{jm}; 0, \tau^{2,s})$ is the normal density function with mean zero and variance $\tau^{2,s}$, evaluated at a_{jm} .

The number of integration points M required to obtain a sufficiently accurate approximation is determined by evaluating the approximation of the target quantity (DIC, AIC) with increasing values of M (7, 11, 17, etc.) and choosing the value of M for which the target quantity changes by less than 0.01 from the previous value. Here the candidate values for M are chosen to increase approximately by 50% while remaining odd so that one of the quadrature locations is at the posterior mean. Furr (2017) finds this approach to be accurate in simulations for linear mixed models where the adaptive quadrature approximation can be compared with the closed form integrals.

Appendix D

Monte Carlo Error for the DIC and WAIC effective number of parameters

For the DIC effective number of parameters, p_D , we can make use of the well-known method for estimating the Monte-Carlo error for the mean of a quantity across MCMC iterations. Let a quantity computed in MCMC iteration s ($s = 1, \dots, S$) be denoted γ_s , so the point estimate of the expectation of γ is

$$\bar{\gamma} = \frac{1}{S} \sum_{s=1}^S \gamma_s.$$

Then the squared Monte Carlo error (or Monte Carlo error variance) is estimated as

$$\text{MCerr}^2(\bar{\gamma}) = \frac{1}{S_{\text{eff}}} \left[\frac{1}{S-1} \sum_{s=1}^S (\gamma_s - \bar{\gamma})^2 \right], \quad (27)$$

where S_{eff} is the effective sample size.

For the effective number of parameter approximation in (6) proposed by Plummer (2008), we can obtain the Monte Carlo error variance by substituting

$$\gamma_s = \frac{1}{2} \log \left\{ \frac{f(\mathbf{y}_s^{\text{r1}} | \boldsymbol{\theta}_s^1)}{f(\mathbf{y}_s^{\text{r1}} | \boldsymbol{\theta}_s^2)} \right\} + \frac{1}{2} \log \left\{ \frac{f(\mathbf{y}_s^{\text{r2}} | \boldsymbol{\theta}_s^2)}{f(\mathbf{y}_s^{\text{r2}} | \boldsymbol{\theta}_s^1)} \right\}$$

in (27).

For the effective number of parameter approximation in (4) proposed by Spiegelhalter et al. (2002), we assume that the variation due to $E_{\boldsymbol{\theta}|\mathbf{y}}[-2 \log f(\mathbf{y}|\boldsymbol{\theta})]$ dominates and use

$$\gamma_s = -2 \log f(\mathbf{y}|\boldsymbol{\theta}_s).$$

For the WAIC effective number of parameters, p_W , we use expressions for the Monte Carlo error of sample variances (see, e.g. White, 2010). Let the variance of γ_s over MCMC

iterations be denoted $v(\gamma)$,

$$v(\gamma) = \frac{1}{S-1} \sum_{s=1}^S (\gamma_s - \bar{\gamma})^2 = \frac{1}{S} \sum_{s=1}^S T_s, \quad T_s = \frac{S}{S-1} (\gamma_s - \bar{\gamma})^2$$

Then the Monte Carlo error variance is estimated as

$$\text{MCerr}^2(v(\gamma)) = \frac{1}{S_{\text{eff}} \times S} \sum_{s=1}^S (T_s - v(\gamma))^2, \quad (28)$$

The conditional version of the effective number of parameters is given by the sum over all units of the posterior variances of the pointwise log posterior densities,

$$p_{\text{Wc}} = \sum_{j=1}^J \sum_{i=1}^{n_j} \text{Var}_{\boldsymbol{\omega}, \boldsymbol{\zeta} | \mathbf{y}} [\log f_c(y_{ij} | \boldsymbol{\omega}, \boldsymbol{\zeta}_j)].$$

The posterior variance $\text{Var}_{\boldsymbol{\omega}, \boldsymbol{\zeta} | \mathbf{y}} [\log f_c(y_{ij} | \boldsymbol{\omega}, \boldsymbol{\zeta}_j)]$ for a given unit is estimated by $v(\gamma_{ij})$ with

$$\gamma_{ijs} = \log f_c(y_{ij} | \boldsymbol{\omega}_s, \boldsymbol{\zeta}_{js}),$$

where we have added subscripts ij to identify the unit, and has Monte Carlo error variance $\text{MCerr}^2(v(\gamma_{ij}))$ given in (28). The variance of the sum of the independent contributions $v(\gamma_{ij})$ to \hat{p}_{Wc} is the sum of the variances of these contributions,

$$\text{MCerr}^2(\hat{p}_{\text{Wc}}) = \sum_{j=1}^J \sum_{i=1}^{n_j} \text{MCerr}^2(v(\gamma_{ij})).$$

For the marginal version of the effective number of parameters, p_{Wm} , we define

$$\gamma_{js} = \log f_m(\mathbf{y}_j | \boldsymbol{\omega}_s, \boldsymbol{\psi}_s)$$

and

$$\text{MCerr}^2(\hat{p}_{\text{Wm}}) = \sum_{j=1}^J \text{MCerr}^2(v(\gamma_j)).$$